

<u>PRINCIPLE</u>	<u>DEFINITION</u>
<u>A) JUSTICE</u>	<u>Fair distribution of AI benefits and risks</u>
<u>B) NON-MALEFICENCE</u>	<u>Ensuring AI does not harm individuals or society</u>
<u>C) AUTONOMY</u>	<u>Respecting users' right to control their data and decisions</u>
<u>D) SUSTAINABILITY</u>	<u>Designing AI to be environmentally friendly</u>

Justice

Example: An AI system used in healthcare is trained on diverse datasets to ensure diagnostic accuracy across populations. It avoids bias that could lead to underdiagnosis in certain ethnic groups.

In your work: NutriTrack could implement fairness-aware algorithms using tools like IBM AI Fairness 360 to ensure food recommendations don't disadvantage rural or low-income users.

Non-maleficence

Example: A social media platform uses content moderation AI that flags harmful misinformation without censoring legitimate information. It protects users from psychological and societal harm.

In MindScape: Ensuring microlearning content doesn't inadvertently reinforce negative mental health stereotypes, and using secure messaging protocols to prevent data breaches, aligns with this principle.

Autonomy

Example: A wellness app lets users control what personal data they share with therapists or coaches. No data is collected without explicit consent, and users can revoke access anytime.

In MindScape: Building user dashboards for permission management and enabling anonymous journaling options respects user autonomy while enhancing trust.

Sustainability

Example: A machine learning model runs on cloud infrastructure optimized for energy efficiency, using green data centers or scheduled downtimes to reduce carbon footprint.
For NutriTrack or AgriSmart: You could design the app to encourage sustainable farming practices through AI-powered tips—minimizing fertilizer overuse or water waste via smart predictions.